

# Locality failover

🕒 4 minute read ✓ page test

---

Configure locality failover

Verify traffic stays in `region1.zone1`

Failover to `region1.zone2`

Failover to `region2.zone3`

Failover to `region3.zone4`

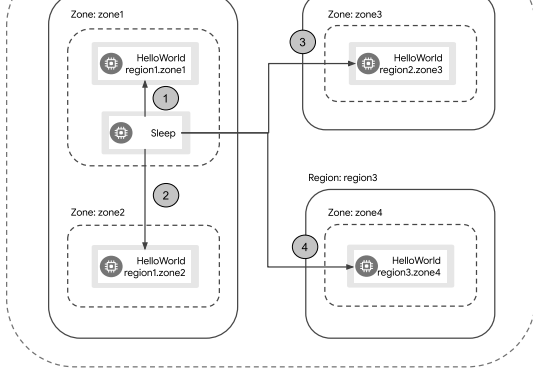
Next steps

Follow this guide to configure your mesh for locality failover.

Before proceeding, be sure to complete the steps under [before you begin](#).

In this task, you will use the `sleep` pod in `region1.zone1` as the source of requests to the `HelloWorld` service. You will then trigger failures that will cause failover between localities in the following sequence:





Locality failover sequence

Internally, Envoy priorities are used to control failover. These

priorities will be assigned as follows for traffic originating from the `sleep pod` (in `region1 zone1`):

Priority	Locality	Details
0	<code>region1.zone1</code>	Region, zone, and sub-zone all match.
1	None	Since this task doesn't use sub-zones, there are no matches for a different sub-zone.
2	<code>region1.zone2</code>	Different zone within the same region.

3	<code>region2.zone3</code>	No match, however failover is defined for <code>region1-&gt;region2</code> .
4	<code>region3.zone4</code>	No match and no failover defined for <code>region1-&gt;region3</code> .

## Configure locality failover

Apply a `DestinationRule` that configures the following:

- Outlier detection for the `HelloWorld` service. This is required in order for failover to function properly. In particular, it

configures the sidecar proxies to know when endpoints for a service are unhealthy, eventually triggering a failover to the next locality.

- **Failover policy** between regions. This ensures that failover beyond a region boundary will behave predictably.
- **Connection Pool policy** that forces each HTTP request to use a new connection. This task utilizes Envoy's drain function to force a failover to the next locality. Once drained, Envoy will reject new connection requests. Since each request uses a new connection, this results in failover immediately following a drain. **This configuration is used for demonstration purposes only.**

```
$ kubectl --context="${CTX_PRIMARY}" apply -n sample -f - <<EOF
apiVersion: networking.istio.io/v1beta1
```

```
kind: DestinationRule
metadata:
  name: helloworld
spec:
  host: helloworld.sample.svc.cluster.local
  trafficPolicy:
    connectionPool:
      http:
        maxRequestsPerConnection: 1
    loadBalancer:
      simple: ROUND_ROBIN
      localityLbSetting:
        enabled: true
        failover:
          - from: region1
            to: region2
    outlierDetection:
      consecutive5xxErrors: 1
      interval: 1s
      baseEjectionTime: 1m
```

EOF

# Verify traffic stays in `region1.zone1`

Call the `HelloWorld` service from the `Sleep` pod:

```
$ kubectl exec --context="${CTX_R1_Z1}" -n sample -c sleep \
  "${(kubectl get pod --context="${CTX_R1_Z1}" -n sample -l \
    app=sleep -o jsonpath='{.items[0].metadata.name}')" \
  -- curl -sSL helloworld.sample:5000/hello
Hello version: region1.zone1, instance: helloworld-region1.zone1-86f77cd7b-
cpxhv
```

Verify that the `version` in the response is `region1.zone`.

Repeat this several times and verify that the response is always the same.



# Failover to region1.zone2

Next, trigger a failover to region1.zone2. To do this, you drain the Envoy sidecar proxy for HelloWorld in region1.zone1:

```
$ kubectl --context="${CTX_R1_Z1}" exec \
  "$(kubectl get pod --context="${CTX_R1_Z1}" -n sample -l app=helloworld \
    -l version=region1.zone1 -o jsonpath='{.items[0].metadata.name}')" \
  -n sample -c istio-proxy -- curl -sSL -X POST 127.0.0.1:15000/drain_liste
ners
```

Call the HelloWorld service from the Sleep pod:

```
$ kubectl exec --context="${CTX_R1_Z1}" -n sample -c sleep \
  "$({kubectl get pod --context="${CTX_R1_Z1}" -n sample -l \
    app=sleep -o jsonpath='{.items[0].metadata.name}')" \
  -- curl -sSL helloworld.sample:5000/hello
Hello version: region1.zone2, instance: helloworld-region1.zone2-86f77cd7b-
cpxhv
```

The first call will fail, which triggers the failover. Repeat the command several more times and verify that the `version` in the response is always `region1.zone2`.

## **Failover to** `region2.zone3`

Now trigger a failover to `region2.zone3`. As you did previously,

configure the HelloWorld in region1.zone2 to fail when called:

```
$ kubectl --context="${CTX_R1_Z2}" exec \
  "$(kubectl get pod --context="${CTX_R1_Z2}" -n sample -l app=helloworld \
  -l version=region1.zone2 -o jsonpath='{.items[0].metadata.name}')" \
  -n sample -c istio-proxy -- curl -sSL -X POST 127.0.0.1:15000/drain_liste
ners
```

Call the HelloWorld service from the Sleep pod:

```
$ kubectl exec --context="${CTX_R1_Z1}" -n sample -c sleep \
  "$(kubectl get pod --context="${CTX_R1_Z1}" -n sample -l \
  app=sleep -o jsonpath='{.items[0].metadata.name}')" \
  -- curl -sSL helloworld.sample:5000/hello
Hello version: region2.zone3, instance: helloworld-region2.zone3-86f77cd7b-
cpxhv
```

The first call will fail, which triggers the failover. Repeat the

command several more times and verify that the `version` in the response is always `region2.zone3`.

## **Failover to** `region3.zone4`

Now trigger a failover to `region3.zone4`. As you did previously, configure the `HelloWorld` in `region2.zone3` to fail when called:

```
$ kubectl --context="${CTX_R2_Z3}" exec \
  "$(kubectl get pod --context="${CTX_R2_Z3}" -n sample -l app=helloworld \
    -l version=region2.zone3 -o jsonpath='{.items[0].metadata.name}')" \
  -n sample -c istio-proxy -- curl -sSL -X POST 127.0.0.1:15000/drain_liste
ners
```

Call the HelloWorld service from the Sleep pod:

```
$ kubectl exec --context="${CTX_R1_Z1}" -n sample -c sleep \
  "$(kubectl get pod --context="${CTX_R1_Z1}" -n sample -l \
    app=sleep -o jsonpath='{.items[0].metadata.name}')" \
  -- curl -sSL helloworld.sample:5000/hello
Hello version: region3.zone4, instance: helloworld-region3.zone4-86f77cd7b-
cpxhv
```

The first call will fail, which triggers the failover. Repeat the command several more times and verify that the `version` in the response is always `region3.zone4`.

**Congratulations!** You successfully configured locality failover!

# Next steps

Cleanup resources and files from this task.