

Mirroring

 4 minute read  page test

This task demonstrates the traffic mirroring capabilities of Istio.

Traffic mirroring, also called shadowing, is a powerful concept that allows feature teams to bring changes to production with as little risk as possible. Mirroring sends a copy of live traffic to a mirrored service. The

mirrored traffic happens out of band of the critical request path for the primary service.

In this task, you will first force all traffic to `v1` of a test service. Then, you will apply a rule to mirror a portion of traffic to `v2`.

Before you begin

- Set up Istio by following the instructions in the [Installation guide](#).

- Start by deploying two versions of the `httpbin` service that have access logging enabled:

httpbin-v1:

```
$ cat <<EOF | istioctl kube-inject -f - | kubectl create -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: httpbin-v1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: httpbin
      version: v1
  template:
    metadata:
```

```
labels:
  app: httpbin
  version: v1
spec:
  containers:
  - image: docker.io/kennethreitz/httpbin
    imagePullPolicy: IfNotPresent
    name: httpbin
    command: ["gunicorn", "--access-logfile", "-", "-b"
, "0.0.0.0:80", "httpbin:app"]
    ports:
    - containerPort: 80
EOF
```

httpbin-v2:

```
$ cat <<EOF | istioctl kube-inject -f - | kubectl create -f
-
apiVersion: apps/v1
```

```
kind: Deployment
metadata:
  name: httpbin-v2
spec:
  replicas: 1
  selector:
    matchLabels:
      app: httpbin
      version: v2
  template:
    metadata:
      labels:
        app: httpbin
        version: v2
    spec:
      containers:
        - image: docker.io/kennethreitz/httpbin
          imagePullPolicy: IfNotPresent
          name: httpbin
          command: ["gunicorn", "--access-logfile", "-", "-b"]
```

```
, "0.0.0.0:80", "httpbin:app"]  
  ports:  
    - containerPort: 80  
EOF
```

httpbin Kubernetes service:

```
$ kubectl create -f - <<EOF
apiVersion: v1
kind: Service
metadata:
  name: httpbin
  labels:
    app: httpbin
spec:
  ports:
    - name: http
      port: 8000
      targetPort: 80
  selector:
    app: httpbin
EOF
```

- Start the `sleep` service so you can use `curl` to provide load:

sleep service:

```
$ cat <<EOF | istioctl kube-inject -f - | kubectl create -f -
-
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sleep
spec:
  replicas: 1
  selector:
    matchLabels:
      app: sleep
  template:
    metadata:
      labels:
        app: sleep
    spec:
      containers:
      - name: sleep
```



```
image: curlimages/curl  
command: ["/bin/sleep","3650d"]  
imagePullPolicy: IfNotPresent
```

EOF

Creating a default routing policy

By default Kubernetes load balances across both versions of the `httpbin` service. In this step, you will change that behavior so that all traffic goes to `v1`.

1. Create a default route rule to route all traffic to v1 of the service:

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: httpbin
spec:
  hosts:
    - httpbin
  http:
    - route:
        - destination:
            host: httpbin
            subset: v1
            weight: 100
    ---
apiVersion: networking.istio.io/v1alpha3
```

```
kind: DestinationRule
metadata:
  name: httpbin
spec:
  host: httpbin
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
EOF
```

Now all traffic goes to the `httpbin:v1` service.

2. Send some traffic to the service:

```
$ export SLEEP_POD=$(kubectl get pod -l app=sleep -o jsonpa
th={.items..metadata.name})
```

```
$ kubectl exec "${SLEEP_POD}" -c sleep -- curl -sS http://httpbin:8000/headers
{
  "headers": {
    "Accept": "*/*",
    "Content-Length": "0",
    "Host": "httpbin:8000",
    "User-Agent": "curl/7.35.0",
    "X-B3-Parentspanid": "57784f8bff90ae0b",
    "X-B3-Sampled": "1",
    "X-B3-Spanid": "3289ae7257c3f159",
    "X-B3-Traceid": "b56eebd279a76f0b57784f8bff90ae0b",
    "X-Envoy-Attempt-Count": "1",
    "X-Forwarded-Client-Cert": "By=spiffe://cluster.local/ns/default/sa/default;Hash=20afebed6da091c850264cc751b8c9306abac02993f80bdb76282237422bd098;Subject=\"\";URI=spiffe://cluster.local/ns/default/sa/default"
  }
}
```

3. Check the logs for `v1` and `v2` of the `httpbin` pods. You should see access log entries for `v1` and none for `v2`:

```
$ export V1_POD=$(kubectl get pod -l app=httpbin,version=v1
-o jsonpath={.items..metadata.name})
$ kubectl logs "$V1_POD" -c httpbin
127.0.0.1 - - [07/Mar/2018:19:02:43 +0000] "GET /headers HTTP/1.1" 200 321 "-" "curl/7.35.0"
```

```
$ export V2_POD=$(kubectl get pod -l app=httpbin,version=v2
-o jsonpath={.items..metadata.name})
$ kubectl logs "$V2_POD" -c httpbin
<none>
```

Mirroring traffic to v2

1. Change the route rule to mirror traffic to v2:

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: httpbin
spec:
  hosts:
    - httpbin
  http:
    - route:
        - destination:
            host: httpbin
            subset: v1
            weight: 100
      mirror:
        host: httpbin
        subset: v2
      mirrorPercentage:
        value: 100.0
EOF
```

This route rule sends 100% of the traffic to `v1`. The last stanza specifies that you want to mirror (i.e., also send) 100% of the same traffic to the `httpbin:v2` service. When traffic gets mirrored, the requests are sent to the mirrored service with their Host/Authority headers appended with `-shadow`. For example, `cluster-1` becomes `cluster-1-shadow`.

Also, it is important to note that these requests are mirrored as “fire and forget”, which means that the responses are discarded.

You can use the `value` field under the `mirrorPercentage` field to mirror a fraction of the

traffic, instead of mirroring all requests. If this field is absent, all traffic will be mirrored.

2. Send in traffic:

```
$ kubectl exec "${SLEEP_POD}" -c sleep -- curl -sS http://httpbin:8000/headers
```

Now, you should see access logging for both `v1` and `v2`. The access logs created in `v2` are the mirrored requests that are actually going to `v1`.

```
$ kubectl logs "$V1_POD" -c httpbin
127.0.0.1 - - [07/Mar/2018:19:02:43 +0000] "GET /headers HTTP/1.1" 200 321 "-" "curl/7.35.0"
127.0.0.1 - - [07/Mar/2018:19:26:44 +0000] "GET /headers HTTP/1.1" 200 321 "-" "curl/7.35.0"
```

```
$ kubectl logs "$V2_POD" -c httpbin
127.0.0.1 - - [07/Mar/2018:19:26:44 +0000] "GET /headers HTTP/1.1" 200 361 "-" "curl/7.35.0"
```

Cleaning up

1. Remove the rules:

```
$ kubectl delete virtualservice httpbin
$ kubectl delete destinationrule httpbin
```

2. Shutdown the `httpbin` service and client:

```
$ kubectl delete deploy httpbin-v1 httpbin-v2 sleep  
$ kubectl delete svc httpbin
```