

Classifying Metrics Based on Request or Response

🕒 6 minute read ✖ page test

It's useful to visualize telemetry based on the type of requests and responses handled by services in your mesh. For example, a bookseller tracks the number of times book reviews are requested. A book review

request has this structure:

```
GET /reviews/{review_id}
```

Counting the number of review requests must account for the unbounded element `review_id`. `GET /reviews/1` followed by `GET /reviews/2` should count as two requests to get reviews.

Istio lets you create classification rules using the `AttributeGen` plugin that groups requests into a fixed number of logical operations. For example, you can create an operation named `GetReviews`, which is a common way to identify operations using the `Open API`

`Spec` `operationId`. This information is injected into request processing as `istio_operationId` attribute with value equal to `GetReviews`. You can use the attribute as a dimension in Istio standard metrics. Similarly, you can track metrics based on other operations like `ListReviews` and `CreateReviews`.

For more information, see the [reference content](#).

Istio uses the Envoy proxy to generate metrics and provides its configuration in the `EnvoyFilter` at `manifests/charts/istio-control/istio-discovery/templates/telemetryv2_1.11.yaml`. As a result, writing classification rules involves adding attributes

to the `EnvoyFilter`.

Classify metrics by request

You can classify requests based on their type, for example `ListReview`, `GetReview`, `CreateReview`.

1. Create a file, for example

`attribute_gen_service.yaml`, and save it with the following contents. This adds the `istio.attributegen` plugin to the `EnvoyFilter`. It also

creates an attribute, `istio_operationId` and populates it with values for the categories to count as metrics.

This configuration is service-specific since request paths are typically service-specific.

```
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: istio-attribute-gen-filter
spec:
  workloadSelector:
    labels:
      app: reviews
  configPatches:
    - applyTo: HTTP_FILTER
      match:
```

context: SIDECAR_INBOUND

proxy:

proxyVersion: '1\9.*'

listener:

filterChain:

filter:

name: "envoy.http_connection_manager"

subFilter:

name: "istio.stats"

patch:

operation: INSERT_BEFORE

value:

name: istio.attribute

typed_config:

"@type": type.googleapis.com/udpa.type.v1.TypedSt

struct

type_url: type.googleapis.com/envoy.extensions.fi

lters.http.wasm.v3.Wasm

value:

config:

```

configuration:
  "@type": type.googleapis.com/google.protobuf.StringValue
  value: |
    {
      "attributes": [
        {
          "output_attribute": "istio_operationId",
          "match": [
            {
              "value": "ListReviews",
              "condition": "request.url_path
== '/reviews' && request.method == 'GET'"
            },
            {
              "value": "GetReview",
              "condition": "request.url_path.
matches('^/reviews/[:alnum:]*$') && request.method == 'GET'"
            }
          ]
        }
      ]
    }

```

```

        },
        {
            "value": "CreateReview",
            "condition": "request.url_path
== '/reviews/' && request.method == 'POST'"
        }
    ]
}
]
}
}
vm_config:
  runtime: envoy.wasm.runtime.null
  code:
    local: { inline_string: "envoy.wasm.attri
butegen" }

```

2. Apply your changes using the following command:


```
$ kubectl -n istio-system apply -f attribute_gen_service.yaml
```

3. Find the `stats-filter-1.11` `EnvoyFilter` resource from the `istio-system` namespace, using the following command:

```
$ kubectl -n istio-system get envoyfilter | grep ^stats-filter-1.11
stats-filter-1.11                2d
```

4. Create a local file system copy of the `EnvoyFilter` configuration, using the following command:

```
$ kubectl -n istio-system get envoyfilter stats-filter-1.11 -o yaml > stats-filter-1.11.yaml
```

5. Open `stats-filter-1.11.yaml` with a text editor and locate the `name: istio.stats` extension configuration. Update it to map `request_operation` dimension in the `requests_total` standard metric to `istio_operationId` attribute. The updated configuration file section should look like the following.

```
name: istio.stats
typed_config:
  '@type': type.googleapis.com/udpa.type.v1.TypedStruct
  type_url: type.googleapis.com/envoy.extensions.filters.ht
tp.wasm.v3.Wasm
  value:
    config:
      configuration:
        "@type": type.googleapis.com/google.protobuf.String
Value
      value: |
        {
          "metrics": [
            {
              "name": "requests_total",
              "dimensions": {
                "request_operation": "istio_operationId"
              }
            }
          ]
        }
```

6. Save `stats-filter-1.11.yaml` and then apply the configuration using the following command:

```
$ kubectl -n istio-system apply -f stats-filter-1.11.yaml
```

7. Add the following configuration to the mesh config. This results in the addition of the `request_operation` as a new dimension to the `istio_requests_total` metric. Without it, a new metric with the name `envoy_request_operation__somevalue__istio_requests_total` is created.

```
meshConfig:
  defaultConfig:
    extraStatTags:
      - request_operation
```

8. Generate metrics by sending traffic to your application.
9. After the changes take effect, visit Prometheus and look for the new or changed dimensions, for example `istio_requests_total`.

Classify metrics by

response

You can classify responses using a similar process as requests. Do note that the `response_code` dimension already exists by default. The example below will change how it is populated.

1. Create a file, for example

`attribute_gen_service.yaml`, and save it with the following contents. This adds the `istio.attributegen` plugin to the `EnvoyFilter` and generates the `istio_responseClass` attribute used by the stats plugin.

This example classifies various responses, such as grouping all response codes in the 200 range as a 2xx dimension.

```
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: istio-attribute-gen-filter
spec:
  workloadSelector:
    labels:
      app: productpage
  configPatches:
    - applyTo: HTTP_FILTER
      match:
        context: SIDECAR_INBOUND
        proxy:
          proxyVersion: '1\9.*'
        listener:
```

```

    filterChain:
      filter:
        name: "envoy.http_connection_manager"
        subFilter:
          name: "istio.stats"
  patch:
    operation: INSERT_BEFORE
    value:
      name: istio.attribute
      typed_config:
        "@type": type.googleapis.com/udpa.type.v1.TypedSt
ruct
        type_url: type.googleapis.com/envoy.extensions.fi
lters.http.wasm.v3.Wasm
        value:
          config:
            configuration:
              "@type": type.googleapis.com/google.protobuf.StringValue
              value: |

```



```
    {
      "attributes": [
        {
          "output_attribute": "istio_response
Class",
          "match": [
            {
              "value": "2xx",
              "condition": "response.code >=
200 && response.code <= 299"
            },
            {
              "value": "3xx",
              "condition": "response.code >=
300 && response.code <= 399"
            },
            {
              "value": "404",
              "condition": "response.code ==
404"
```

429"

```
},  
{  
  "value": "429",  
  "condition": "response.code ==
```

503"

```
},  
{  
  "value": "503",  
  "condition": "response.code ==
```

500 && response.code <= 599"

```
},  
{  
  "value": "5xx",  
  "condition": "response.code >=
```

400 && response.code <= 499"

```
},  
{  
  "value": "4xx",  
  "condition": "response.code >=
```

```
    }  
  ]  
}  
]  
}  
vm_config:  
  runtime: envoy.wasm.runtime.null  
  code:  
    local: { inline_string: "envoy.wasm.attri  
butegen" }
```

2. Apply your changes using the following command:

```
$ kubectl -n istio-system apply -f attribute_gen_service.yaml
```

3. Find the `stats-filter-1.11` EnvoyFilter resource from the `istio-system` namespace, using the

following command:

```
$ kubectl -n istio-system get envoyfilter | grep ^stats-filter-1.11
stats-filter-1.11                2d
```

4. Create a local file system copy of the `EnvoyFilter` configuration, using the following command:

```
$ kubectl -n istio-system get envoyfilter stats-filter-1.11 -o yaml > stats-filter-1.11.yaml
```

5. Open `stats-filter-1.11.yaml` with a text editor and locate the `name: istio.stats` extension configuration. Update it to map `response_code` dimension in the `requests_total` standard metric to

`istio_responseClass` attribute. The updated configuration file section should look like the following.

```
name: istio.stats
typed_config:
  '@type': type.googleapis.com/udpa.type.v1.TypedStruct
  type_url: type.googleapis.com/envoy.extensions.filters.ht
tp.wasm.v3.Wasm
  value:
    config:
      configuration:
        "@type": type.googleapis.com/google.protobuf.String
Value
    value: |
      {
        "metrics": [
          {
            "name": "requests_total",
            "dimensions": {
              "response_code": "istio_responseClass"
            }
          }
        ]
      }
    }
```

6. Save `stats-filter-1.11.yaml` and then apply the configuration using the following command:

```
$ kubectl -n istio-system apply -f stats-filter-1.11.yaml
```

Verify the results

1. Generate metrics by sending traffic to your application.
2. Visit Prometheus and look for the new or changed dimensions, for example `2xx`. Alternatively, use

the following command to verify that Istio generates the data for your new dimension:

```
$ kubectl exec pod-name -c istio-proxy -- curl -sS 'localhost:15000/stats/prometheus' | grep istio_
```

In the output, locate the metric (e.g. `istio_requests_total`) and verify the presence of the new or changed dimension.

Troubleshooting

If classification does not occur as expected, check the following potential causes and resolutions.

Review the Envoy proxy logs for the pod that has the service on which you applied the configuration change. Check that there are no errors reported by the service in the Envoy proxy logs on the pod, (`pod-name`), where you configured classification by using the following command:

```
$ kubectl logs pod-name -c istio-proxy | grep -e "Config Error" -e "envoy wasm"
```

Additionally, ensure that there are no Envoy proxy

crashes by looking for signs of restarts in the output of the following command:

```
$ kubectl get pods pod-name
```