

Enabling Rate Limits using Envoy

🕒 6 minute read ✖ page test

This task shows you how to use Envoy's native rate limiting to dynamically limit the traffic to an Istio service. In this task, you will apply a global rate-limit for the `productpage` service through ingress gateway that allows 1 requests per minute across all instances of the service. Additionally, you will apply a local rate-limit for each

individual `productpage` instance that will allow 10 requests per minute. In this way, you will ensure that the `productpage` service handles a maximum of 1 request per minute through the ingress gateway, but each `productpage` instance can handle up to 10 requests per minute, allowing for any in-mesh traffic.

Before you begin

1. Setup Istio in a Kubernetes cluster by following the instructions in the [Installation Guide](#).
2. Deploy the `Bookinfo` sample application.

Rate limits

Envoy supports two kinds of rate limiting: global and local. Global rate limiting uses a global gRPC rate limiting service to provide rate limiting for the entire mesh. Local rate limiting is used to limit the rate of requests per service instance. Local rate limiting can be used in conjunction with global rate limiting to reduce load on the global rate limiting service.

In this task you will configure Envoy to rate limit traffic to a specific path of a service using both global and local rate limits.

Global rate limit

Envoy can be used to set up global rate limits

for your mesh. Global rate limiting in Envoy uses a gRPC API for requesting quota from a rate limiting service. A reference implementation of the API, written in Go with a Redis backend, is used below.

1. Use the following configmap to configure the reference implementation to rate limit requests to the path `/productpage` at 1 req/min and all other requests at 100 req/min.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: ratelimit-config
data:
  config.yaml: |
    domain: productpage-ratelimit
    descriptors:
      - key: PATH
        value: "/productpage"
        rate_limit:
          unit: minute
          requests_per_unit: 1
      - key: PATH
        rate_limit:
          unit: minute
          requests_per_unit: 100
```

2. Create a global rate limit service which implements Envoy's rate limit service protocol. As a reference, a demo configuration can be found [here](#), which is based on a reference implementation provided by Envoy.
3. Apply an `EnvoyFilter` to the `ingressgateway` to enable global rate

limiting using Envoy's global rate limit filter.

The first patch inserts the `envoy.filters.http.ratelimit` global envoy filter into the `HTTP_FILTER` chain.

The `rate_limit_service` field specifies the external rate limit service, `rate_limit_cluster` in this case.

The second patch defines the `rate_limit_cluster`, which provides the endpoint location of the external rate limit service.

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: filter-ratelimit
  namespace: istio-system
spec:
  workloadSelector:
    # select by label in the same namespace
    labels:
      istio: ingressgateway
```

```

configPatches:
  # The Envoy config you want to modify
  - applyTo: HTTP_FILTER
    match:
      context: GATEWAY
      listener:
        filterChain:
          filter:
            name: "envoy.filters.network.http
_connection_manager"
            subFilter:
              name: "envoy.filters.http.route
r"
        patch:
          operation: INSERT_BEFORE
          # Adds the Envoy Rate Limit Filter in H
TTP filter chain.
          value:
            name: envoy.filters.http.ratelimit
            typed_config:
              "@type": type.googleapis.com/envoy.
extensions.filters.http.ratelimit.v3.RateLimit
              # domain can be anything! Match it
to the ratelimiter service config
              domain: productpage-ratelimit
              failure_mode_deny: true
              timeout: 10s
              rate_limit_service:
                grpc_service:
                  envoy_grpc:
                    cluster_name: rate_limit_clus

```

```

ter
    transport_api_version: V3
  - applyTo: CLUSTER
    match:
      cluster:
        service: ratelimit.default.svc.cluster.local
    patch:
      operation: ADD
      # Adds the rate limit service cluster for rate limit service defined in step 1.
      value:
        name: rate_limit_cluster
        type: STRICT_DNS
        connect_timeout: 10s
        lb_policy: ROUND_ROBIN
        http2_protocol_options: {}
        load_assignment:
          cluster_name: rate_limit_cluster
          endpoints:
            - lb_endpoints:
                - endpoint:
                    address:
                      socket_address:
                        address: ratelimit.default.svc.cluster.local
                        port_value: 8081
EOF

```

4. Apply another `EnvoyFilter` to the

ingressgateway that defines the route configuration on which to rate limit. This adds rate limit actions for any route from a virtual host named *.80.

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: filter-ratelimit-svc
  namespace: istio-system
spec:
  workloadSelector:
    labels:
      istio: ingressgateway
  configPatches:
    - applyTo: VIRTUAL_HOST
      match:
        context: GATEWAY
        routeConfiguration:
          vhost:
            name: ""
            route:
              action: ANY
      patch:
        operation: MERGE
        # Applies the rate limit rules.
        value:
          rate_limits:
            - actions: # any actions in here
              - request_headers:
                  header_name: ":path"
                  descriptor_key: "PATH"
```

EOF

Local rate limit

Envoy supports local rate limiting of L4 connections and HTTP requests. This allows you to apply rate limits at the instance level, in the proxy itself, without calling any other service.

The following `EnvoyFilter` enables local rate limiting for any traffic through the `productpage` service. The `HTTP_FILTER` patch inserts the `envoy.filters.http.local_ratelimit` local envoy filter into the HTTP connection manager filter chain. The local rate limit filter's `token bucket` is configured to allow 10 requests/min. The filter is also configured to add an `x-local-rate-limit` response header to requests that are blocked.

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
```

```
kind: EnvoyFilter
metadata:
  name: filter-local-ratelimit-svc
  namespace: istio-system
spec:
  workloadSelector:
    labels:
      app: productpage
  configPatches:
    - applyTo: HTTP_FILTER
      match:
        context: SIDECAR_INBOUND
        listener:
          filterChain:
            filter:
              name: "envoy.filters.network.http_connection_manager"
            patch:
              operation: INSERT_BEFORE
              value:
                name: envoy.filters.http.local_ratelimit
                typed_config:
                  "@type": type.googleapis.com/udpa.type.v1.TypedStruct
                  type_url: type.googleapis.com/envoy.extensions.filters.http.local_ratelimit.v3.LocalRateLimit
                  value:
                    stat_prefix: http_local_rate_limiter
                    token_bucket:
                      max_tokens: 10
```

ed

ced

EOF

```
tokens_per_fill: 10
fill_interval: 60s
filter_enabled:
  runtime_key: local_rate_limit_enabl

  default_value:
    numerator: 100
    denominator: HUNDRED
filter_enforced:
  runtime_key: local_rate_limit_enfor

  default_value:
    numerator: 100
    denominator: HUNDRED
response_headers_to_add:
- append: false
  header:
    key: x-local-rate-limit
    value: 'true'
```

The above configuration applies local rate limiting to all vhosts/routes. Alternatively, you can restrict it to a specific route.

The following `EnvoyFilter` enables local rate limiting for any traffic to port 80 of the

productpage service. Unlike the previous configuration, there is no `token_bucket` included in the `HTTP_FILTER` patch. The `token_bucket` is instead defined in the second (`HTTP_ROUTE`) patch which includes a `typed_per_filter_config` for the `envoy.filters.http.local_ratelimit` local envoy filter, for routes to virtual host `inbound|http|9080`.

```
$ kubectl apply -f - <<EOF
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: filter-local-ratelimit-svc
  namespace: istio-system
spec:
  workloadSelector:
    labels:
      app: productpage
  configPatches:
    - applyTo: HTTP_FILTER
      match:
        context: SIDECAR_INBOUND
        listener:
          filterChain:
```

```

        filter:
            name: "envoy.filters.network.http_connection_manager"
        patch:
            operation: INSERT_BEFORE
            value:
                name: envoy.filters.http.local_ratelimit
                typed_config:
                    "@type": type.googleapis.com/udpa.type.v1.TypedStruct
                    type_url: type.googleapis.com/envoy.extensions.filters.http.local_ratelimit.v3.LocalRateLimit
                    value:
                        stat_prefix: http_local_rate_limiter
- applyTo: HTTP_ROUTE
        match:
            context: SIDECAR_INBOUND
            routeConfiguration:
                vhost:
                    name: "inbound|http|9080"
                    route:
                        action: ANY
        patch:
            operation: MERGE
            value:
                typed_per_filter_config:
                    envoy.filters.http.local_ratelimit:
                        "@type": type.googleapis.com/udpa.type.v1.TypedStruct
                        type_url: type.googleapis.com/envoy.e

```

xtensions.filters.http.local_ratelimit.v3.LocalRate
Limit

value:

stat_prefix: http_local_rate_limite

r

token_bucket:

max_tokens: 10

tokens_per_fill: 10

fill_interval: 60s

filter_enabled:

runtime_key: local_rate_limit_ena

bled

default_value:

numerator: 100

denominator: HUNDRED

filter_enforced:

runtime_key: local_rate_limit_enf

orced

default_value:

numerator: 100

denominator: HUNDRED

response_headers_to_add:

- append: false

header:

key: x-local-rate-limit

value: 'true'


EOF

Verify the results

Verify global rate limit

Send traffic to the Bookinfo sample. Visit `http://$GATEWAY_URL/productpage` in your web browser or issue the following command:

```
$ curl "http://$GATEWAY_URL/productpage"
```

 `$GATEWAY_URL` is the value set in the Bookinfo **example**.

You will see the first request go through but every following request within a minute will get a 429 response.

Verify local rate limit

Although the global rate limit at the ingress gateway limits requests to the `productpage` service at 1 req/min, the local rate limit for `productpage` instances allows 10 req/min. To confirm this, send internal `productpage` requests, from the `ratings` pod, using the following `curl` command:

```
$ kubectl exec "$(kubectl get pod -l app=ratings -o jsonpath='{.items[0].metadata.name}')" -c ratings -- curl -sS productpage:9080/productpage | grep -o "<title>.*</title>"  
<title>Simple Bookstore App</title>
```

You should see no more than 10 req/min go through per `productpage` instance.